



# Modul

## - Internet of Things (IoT) -

**04-Vorlesung**


Prof. Dr. Marcel Tilly

Fakultät Informatik, Cloud Computing



21. März	Einführung in das Internet der Dinge
28. März	IoT Architekturen
4. April	Things und Sensoren
11. April	From Device to Cloud
18. April	Vorlesungsfrei – Ostern
25. April	IoT Analytics
02. Mai	Big Data in IoT
9. Mai	Data Exploration
16. Mai	IoT Plattformen
23. Mai	Entwicklung einer IoT Lösung
30. Mai	Vorlesungsfrei; Christi Himmelfahrt
05. Juni	opt. Gastvortrag – Digitalisierung
13. Juni	Data Science in IoT
20. Juni	Vorlesungsfrei – Fronleichnam
27. Juni	Intelligente Cloud und intelligente Edge
04. Juli	PStA Abschlusspraesentationen

PStA



# What is Raspberry Pi ?



- Arguably the most popular single board computer (SBC)
  - Easy to get started with because basically every problem is documented
- Add a computer with a OS to practically anything
  - NOTE: Do not expect it to perform as well as your laptop
- Support for a vast array of peripherals (thanks to the Linux kernel)
  - USB devices, networking, displays, cameras, audio etc.



# What can you do with a Pi? (not so cool)



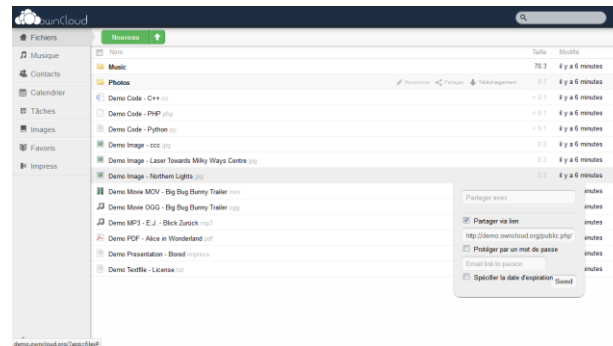
<https://www.raspberrypi.org/magpi/magic-mirror/>



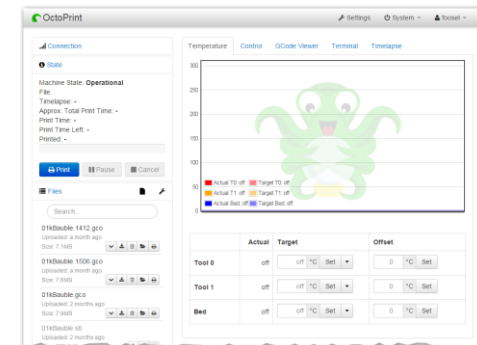
<http://mymediaexperience.com/raspberry-pi-xbmc-with-raspbmc/>



<https://learn.adafruit.com/pigrrl-2/overview>




<https://vadelmapii.com/blogi/yllapida-omaa-dropbox-kloonias-raspberry-pilla-kayttaen-owncloudia>



<https://github.com/foosel/OctoPrint>


# What can you do with a Pi? (cool)



<https://hackaday.io/project/1279-ramanpi-raman-spectrometer>




<https://www.raspberrypi.org/blog/real-time-depth-perception-with-the-compute-module/>




<https://www.raspberrypi.org/blog/real-time-depth-perception-with-the-compute-module/>



<https://hackaday.io/project/1269-mashberry-beer-brewing-with-raspberry-pi>



<http://www.zdnet.com/article/build-your-own-supercomputer-out-of-raspberry-pi-boards/>



<https://www.raspberrypi.org/magpi/fabsca-n-pi-project-3d-scanning-for-all/>



# Raspberry Pi models



**B Rev 1**

0002



**B Rev 1 links**

0003



**A**

0008



**B Rev 2 (256 MB)**

0004



**B Rev 2 (China)**

000f



**B Rev 2.1 (UK)**

000e



**B Rev 2 (Chinese)**

000d



**B Rev 2 (Blue Pi)**

000d



**Compute  
Module**

000d



**B+**

0010



**B+ (Chinese)**

0010



**A+**

0012



**2B**

a01041



**Zero**

900092



**3B**

a02082





**Raspberry Pi®  
family**

Feb 29th 2016

**RasPi.TV**

# Plattform: Raspberry Pi



	Raspberry Pi 3 Model B	Raspberry Pi Zero	Raspberry Pi 2 Model B	Raspberry Pi Model B+
Introduction Date	2/29/2016	11/25/2015	2/2/2015	7/14/2014
SoC	BCM2837	BCM2835	BCM2836	BCM2835
CPU	Quad Cortex A53 @ 1.2GHz	ARM11 @ 1GHz	Quad Cortex A7 @ 900MHz	ARM11 @ 700MHz
Instruction set	ARMv8-A	ARMv6	ARMv7-A	ARMv6
GPU	400MHz VideoCore IV	250MHz VideoCore IV	250MHz VideoCore IV	250MHz VideoCore IV
RAM	1GB SDRAM	512 MB SDRAM	1GB SDRAM	512MB SDRAM
Storage	micro-SD	micro-SD	micro-SD	micro-SD
Ethernet	10/100	none	10/100	10/100
Wireless	802.11n / Bluetooth 4.0	none	none	none
Video Output	HDMI / Composite	HDMI / Composite	HDMI / Composite	HDMI / Composite
Audio Output	HDMI / Headphone	HDMI	HDMI / Headphone	HDMI / Headphone
GPIO	40	40	40	40
Price	\$35	\$5	\$35	\$35

# Raspberry Pi advantages and disadvantages



- Cheap (price per performance)
- Well documented
- Availability
  - Also in terms of add-ons (HATs)
- Versatile
- Compact (especially Zero)



- Scary linux (learning required)
- Not real time\*
- No ADC (easy to add though)
- PWM possible but limited frequency

\* It is possible to install a RTOS on the Pi



# OS options



NOOBS



RASPBIAN



WINDOWS 10 IOT CORE



RISC OS

A non-Linux distribution

with Pixel or Lite



OSMC



LIBREELEC



UBUNTU MATE



SNAPPY UBUNTU CORE

# Installation process

Get SD card (micro on newer Pi's)


Format to FAT32 (for example with SD card formatter)

Using NOOBS:

- Download NOOBS (sd-cards with pre-installed NOOBS can purchased)
- Unzip and copy all contents to SD card (takes a while)
- Boot

Using a disk image writer:

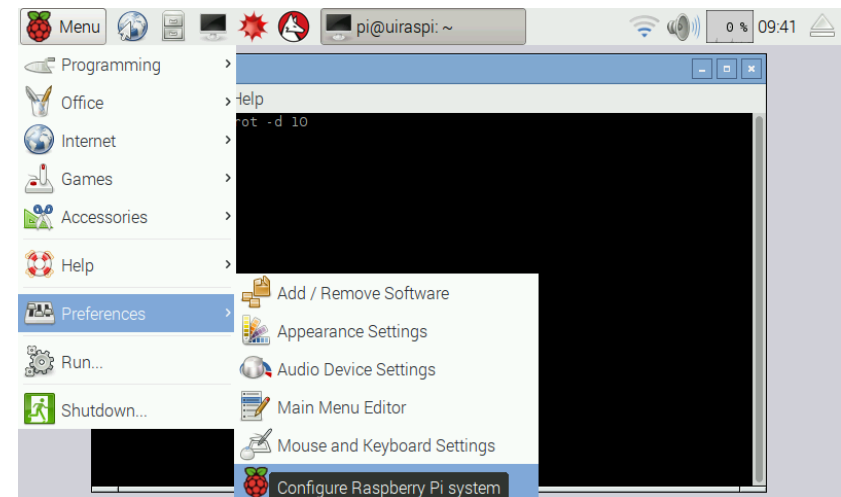
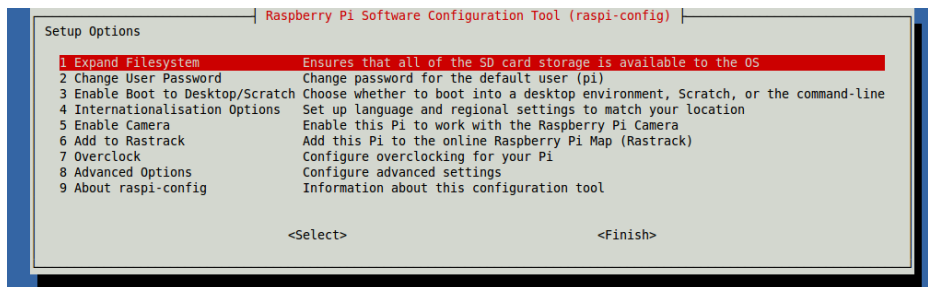
- Download the disk image (.img)
- Write the disk image on the SD card ("hard drive")
  - Windows: Win32 disk image writer or Etcher
  - MacOSX / Linux: Use dd (or image writer with GUI such as Etcher)
- Boot



## On first boot...




- NOOBS installer has GUI (Raspbian recommended)
- Boots into raspi-config (you can run it with “sudo raspi-config”)
- Expand file system, change password and change keyboard layout, enable ssh etc.



# Terminal or Pixel



There are two options when booting: shell or Pixel desktop (graphical session)



```
Debian GNU/Linux wheezy/sid raspberrypi tty1
raspberrypi login: pi
Password:
Last login: Tue Aug 21 21:24:50 EDT 2012 on tty1
Linux raspberrypi 3.1.9+ #168 PREEMPT Sat Jul 14 18:56:31 BST 2012 armv6l

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.

Type 'startx' to launch a graphical session

pi@raspberrypi ~ $
```




# Connect to WLAN



## In GUI:

- Choose WiFi network and connect



## Without GUI :

- `sudo iwlist wlan0 scan`
- `sudo nano /etc/wpa_supplicant/wpa_supplicant.conf`  
configuration
- Add this to the bottom of the file

`// scan for networks`

`// open wpa_supplicant`

```
network={
    ssid="aalto open"
    proto=RSN
    key_mgmt=NONE
}
```

# Data Storage



## Challenges

**Amount of Data (Big Data)** - Volume

**Speed of Data (Data Rate)** - Velocity

**Data format (semantics)** - Variety

Data trust

**Data quality**

DSVGO (privacy)

Encryption (Transport security)

Compression

Data Selection

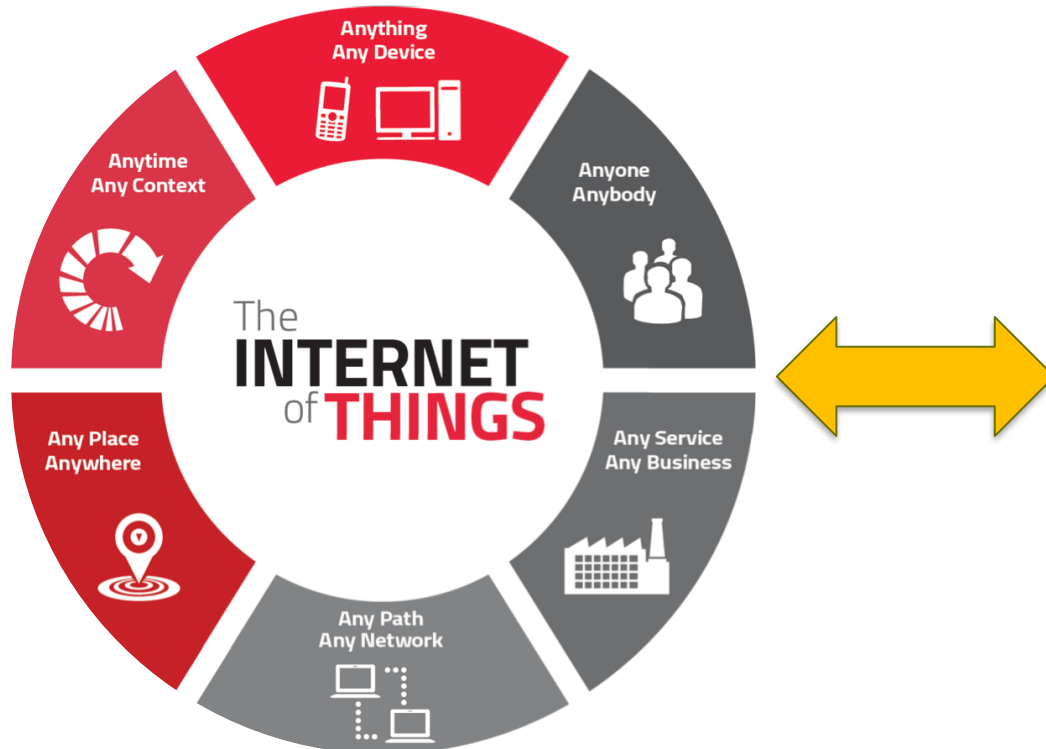
## Technology

SQL

Map-and-Reduce (Hadoop)

Stream Processing (Spark)

# IoT Meets Big Data



# When does data become Big Data

Big data is generally defined by its “three Vs”: volume, velocity, and variety. Thinkers in the field have occasionally argued for adding additional “Vs” (such as veracity and value), but the classic three Vs provide a nice overview of what defines big data.

- **Volume** is the “big” in big data. Current volumes of data (4 zettabytes worldwide this year and growing) dwarf anything we have had to tackle in human history.
- **Velocity** refers to the rate at which big data needs to be processed. This is both a factor of how quickly new data gets created, but also that we often want to look at live data streams for fresh data, like Twitter feeds for sentiment analysis or real-time Internet of Things (IoT) sensor data streams.
- **Variety** refers to unstructured (images, audio, and the like) or semi-structured (JSON documents).

# Big Data Value Chain



**Collection** – Structured, unstructured and semi-structured data from multiple sources

**Ingestion** – loading vast amounts of data onto a single data store

**Discovery & Cleansing** – understanding format and content; clean up and formatting

**Integration** – linking, entity extraction, entity resolution, indexing and data fusion

**Analysis** – Intelligence, statistics, predictive and text analytics, machine learning

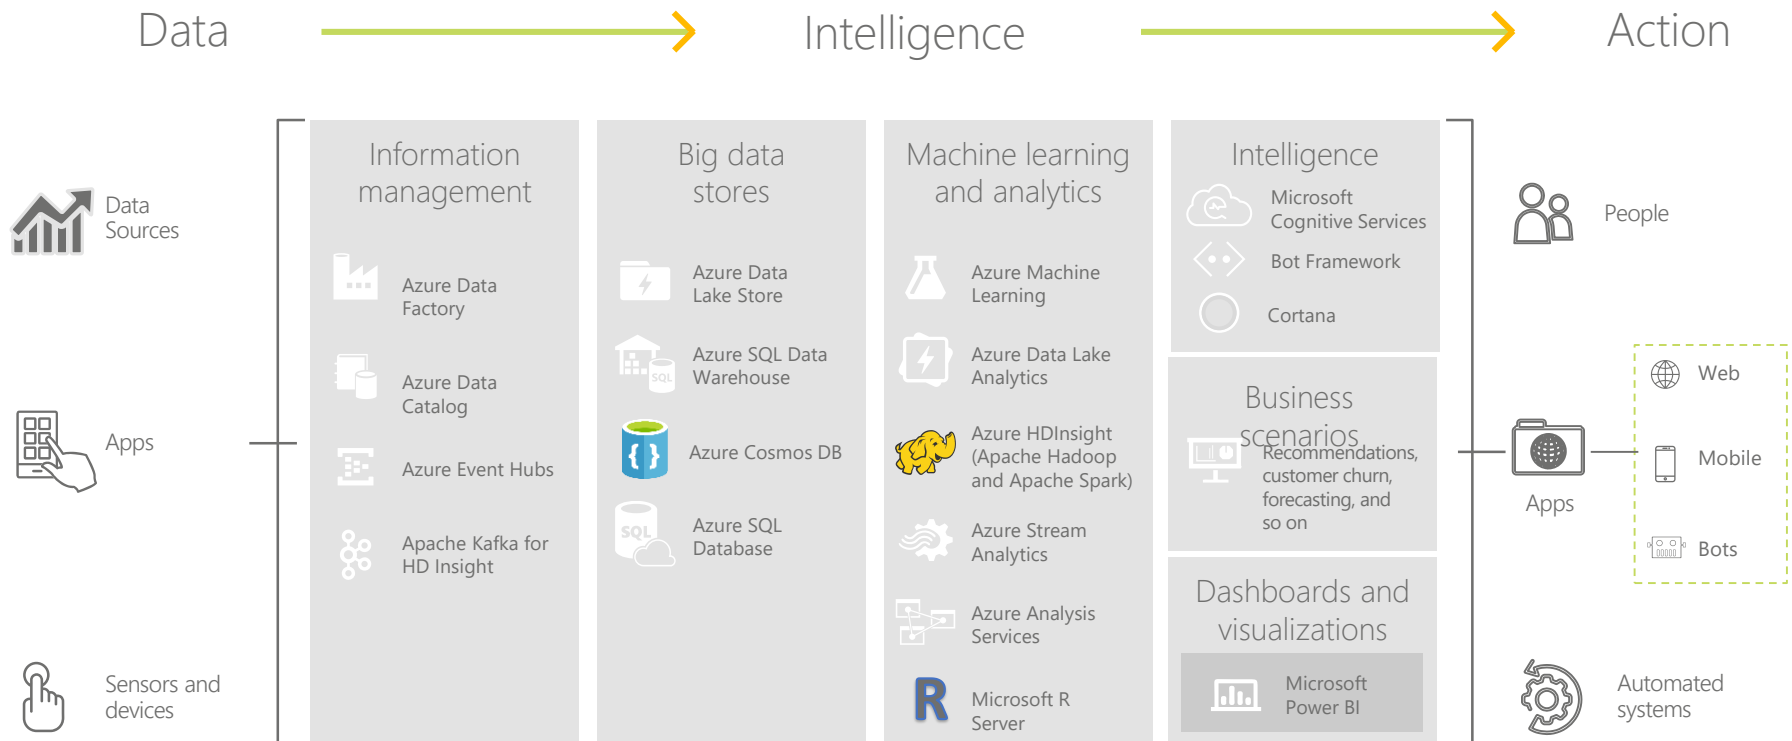
**Delivery** – querying, visualization, real time delivery on enterprise-class availability

Source O'Reilly Strata 2012

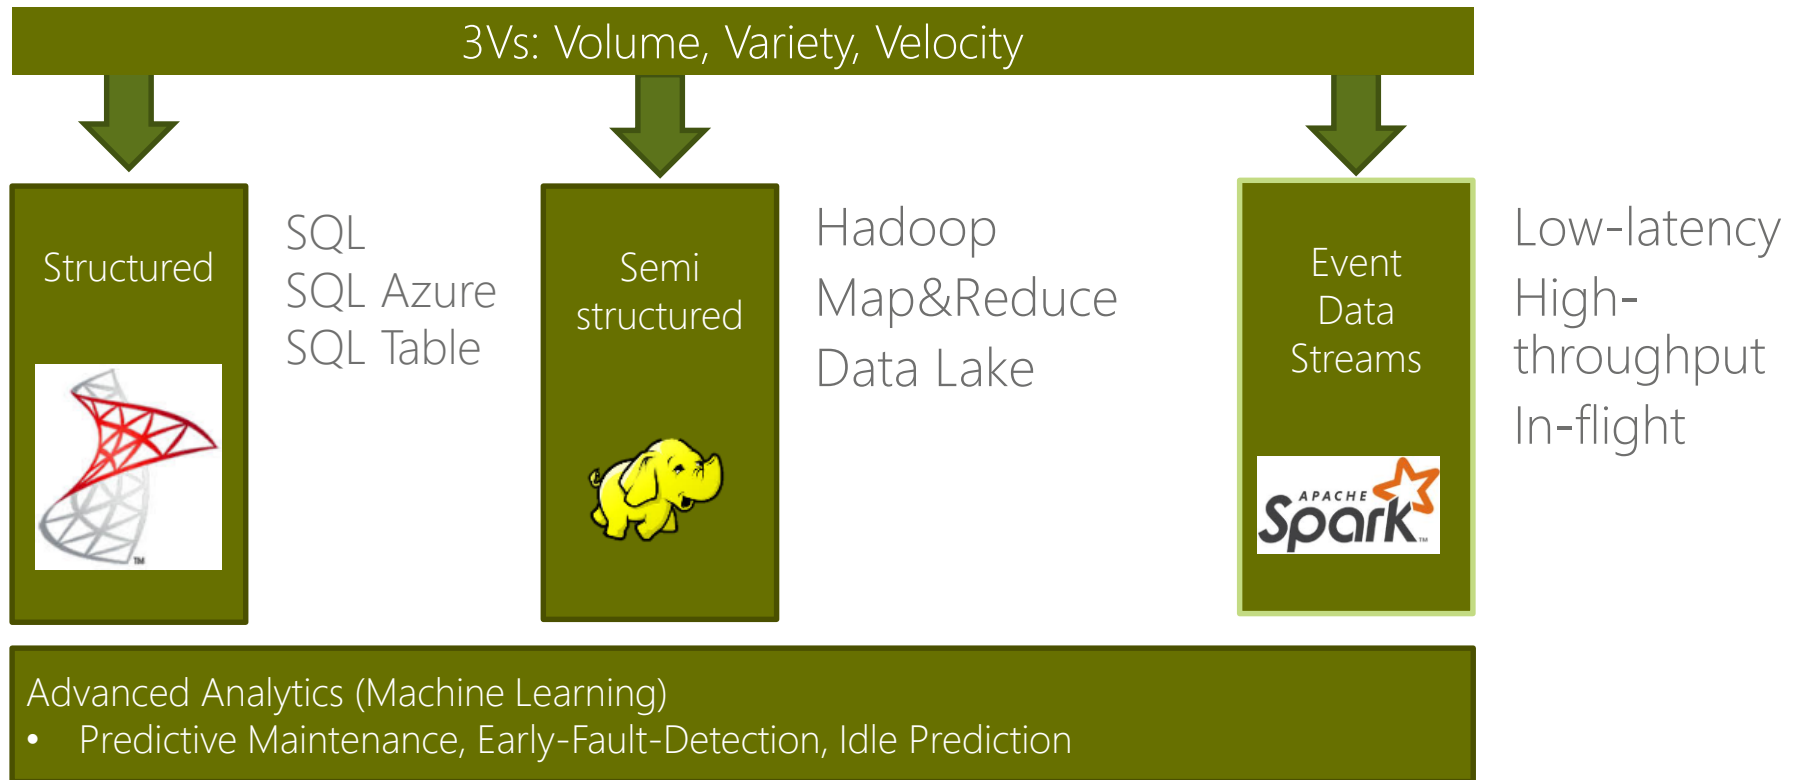


# Azure Big Data platform


## Transform data into intelligent actions and predictions



# From(data).To(insights)




# Understanding Streaming Data (1)



- **Question:** “How many red cars are in the parking lot?”
- Prerequisite: Data is static! (no cars leaving or entering!)
- Answering with a **relational database:**
  - Walk out to the parking lot
  - Count vehicles that are: ‘Red’ and ‘Cars’

```
(FROM vhc1 in ParkingLot
WHERE vhc1.type  = “CAR”
&&    vhd1.color = “RED”
SELECT vhc1).Count()
```

# Understanding Streaming Data (2)



“How many **trucks (red cars)** went from right to left in the last **10 seconds**”?

“How many **green cars (= electrical cars)** are passing every **60 minutes**”?


“What is the **running average speed** of cars on the left lane **compared** to the right lane?”

This is the streaming data paradigm in a nutshell –  
ask questions about data in flight.

# Data Stream in a Nutshell



- A data stream is a continuous sequence of data tuples
  - Think of standard tuples of relational databases
  - + time information (**timestamp = TS**)



- **That means:**
  - **Data is moving!**
  - Continuously generated (assumed **infinite!**)
- Potentially **high pace**.
- System has to process data without first storing everything







- Sensor networks for environmental monitoring
  - Avalanche risk level computation
  - Insights for agriculture
  - Air pollution (urban) monitoring
- Real-time analysis of stock market changes
  - Computing statistics over streams, e.g., for decision support
  - Opportunities for reacting in real-time
  - Even with fully automated means: algorithmic trading
- Social Media Analysis
  - Sentiment analysis of products
  - News and trend analysis
- Industrial monitoring
  - Energy consumption and other key performance indicators (KPIs)
  - Predictive maintenance

# Streaming Data Paradigm

	Database Applications	Data Streaming Applications
<b>Concept</b>	Persist data (&relations) (Data at Rest)	Volatile data streams (Data in Motion)
<b>Strategy</b>	Random access	Sequential access
<b>Query Paradigm</b>	Ad-hoc queries or requests	Continuous standing queries
<b>Memory/Storage</b>	(in theory) unlimited storage	Main memory limitations
<b>Execution Time</b>	Little or no time requirements  (Seconds, hours, days)	Consideration of the order of the input ( $< 1$ sec.)
<b>Data Rate</b>	Low rate (Hundreds of events/sec)	High rate (tens of thousands of events/sec or more)
<b>Query Semantics</b>	Declarative relational analytics	Declarative relational <i>and temporal</i> analytics
<b>Accuracy</b>	Assumes exact data	Assumes inaccurate data



# Data Stream Terminology

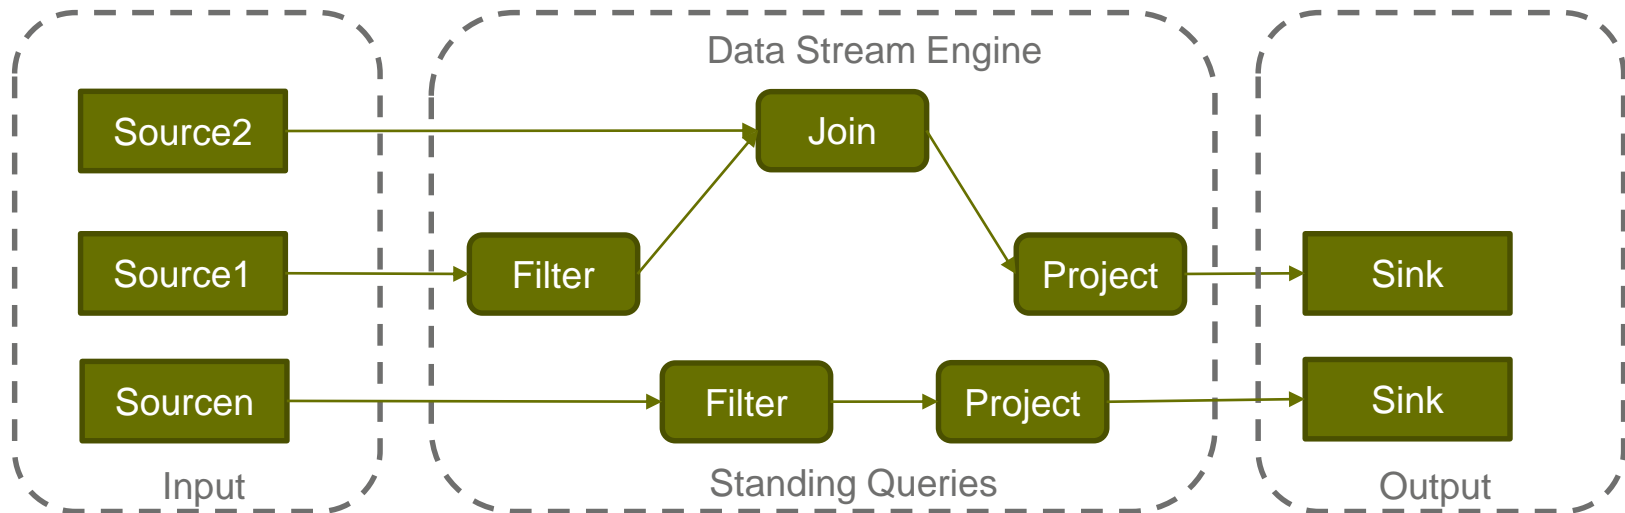


Term	Notation	Description	Technology Sample
Tuple/ payload	$p = \langle k1:v1, k2:v2, \dots \rangle$	A un-ordered list of named elements (e.g. JSON)	
Event	$e = \langle ts, p \rangle$	A timestamped tuple	
Stream	$s$	A stream is a possibly infinite sequence of events	
Source	$src \mid \rangle s$	A source is emitting continuously events	Storm: Spout StreamInsight: IObservable
Operator	$S_{1,in} \dots S_{n,in} \mid \rangle o \mid \rangle s_{1,out} \dots s_{m,out}$	Processor of $n$ input streams producing $m$ output stream of events	Storm: Bolt StreamInsight: Function
Topology	$G = (src[1..n], o, sk)$	A graph of calculation represented as a network (= query) with $n$ sources and 1 sink	
Sink	$s \mid \rangle sk$	Consumes results of a Stream	StreamInsight: IObservable

# Data Stream Topology

A standing query is an *instance* of a topology (=query graph)

A *data stream engine* handles heavy lifting for data streams



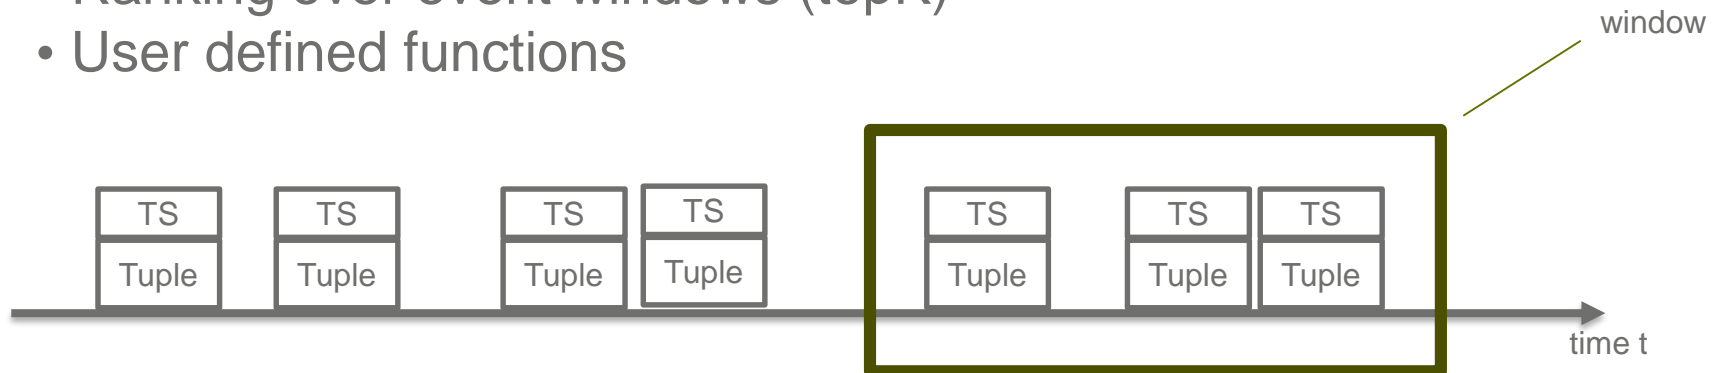


## •Stream Operations

- Selection of events = filter
- Calculations on the tuple payload = projection
- Correlation of streams = join
- Stream partitioning = grouping

## •Window Operations

- Aggregation (sum, count, min, max...) over event windows
- Ranking over event windows (topK)
- User defined functions





# Query Expressiveness

Filter  
Projection

```
var result = from e in inputStream
              where e.id > 3
              select new {
                  id = e.id,
                  W = (double)e.intW / 10 };
```

Correlation (Join)

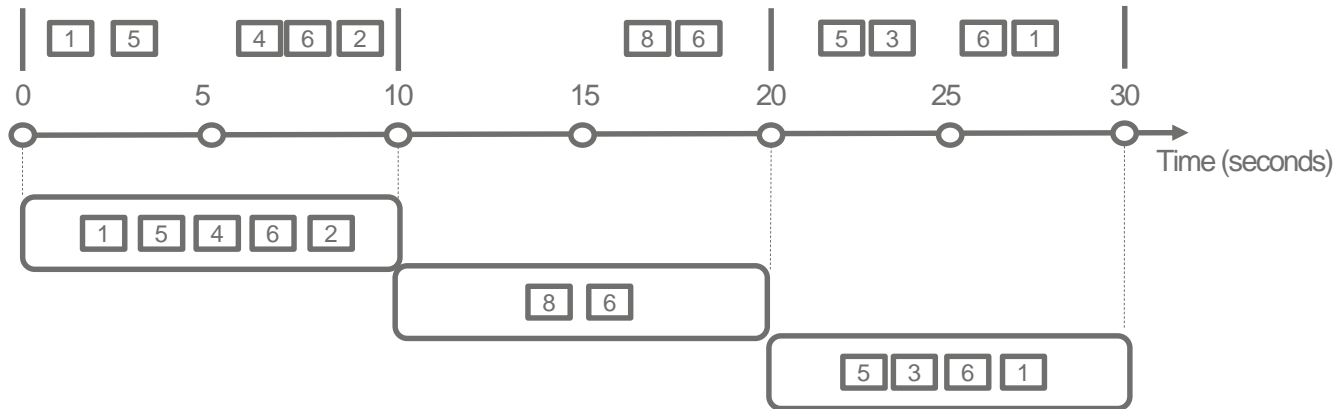
```
var result = from eLeft in inputStream1
              join eRight in inputStream2
              on eLeft.id equals eRight.id
              select new {
                  id = eLeft.id,
                  diff = eLeft.W - eRight.w
              };
```

Projection

# Time Windows



A 10-second tumbling window



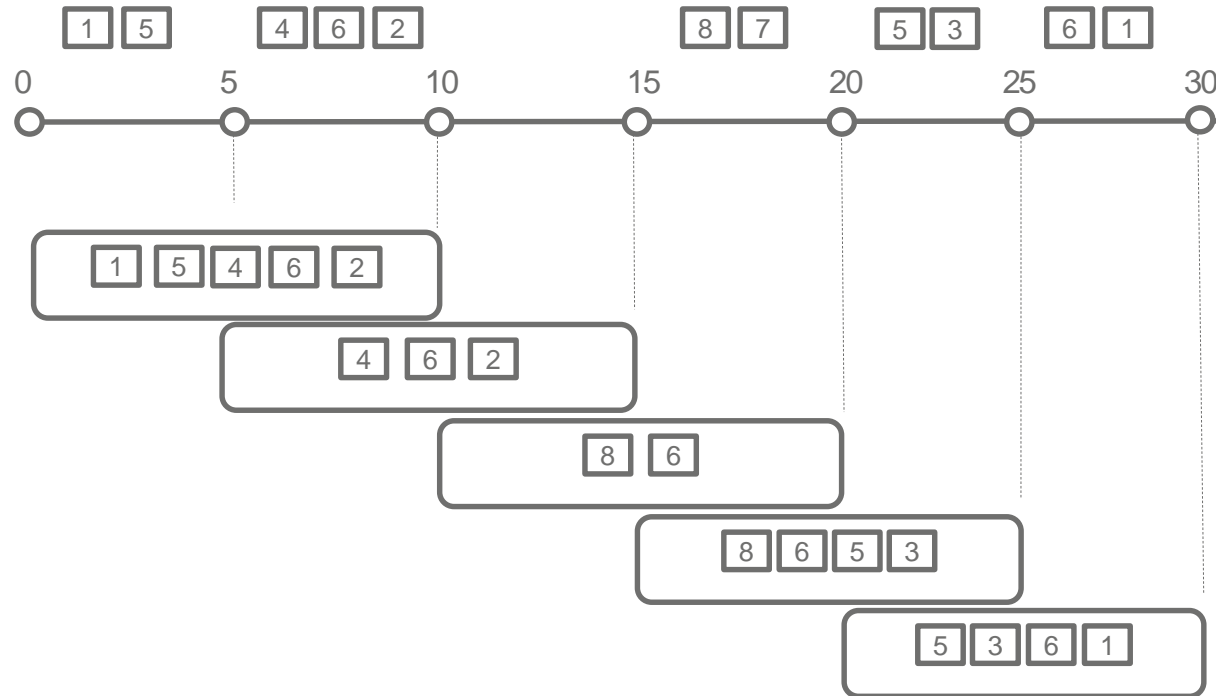
How many vehicles entered each toll both every 10 seconds?

```
SELECT TollId, COUNT(*) FROM EntryStream  
TIMESTAMP BY EntryTime  
GROUP BY TollId, TumblingWindow(second,10)
```

# Time Windows

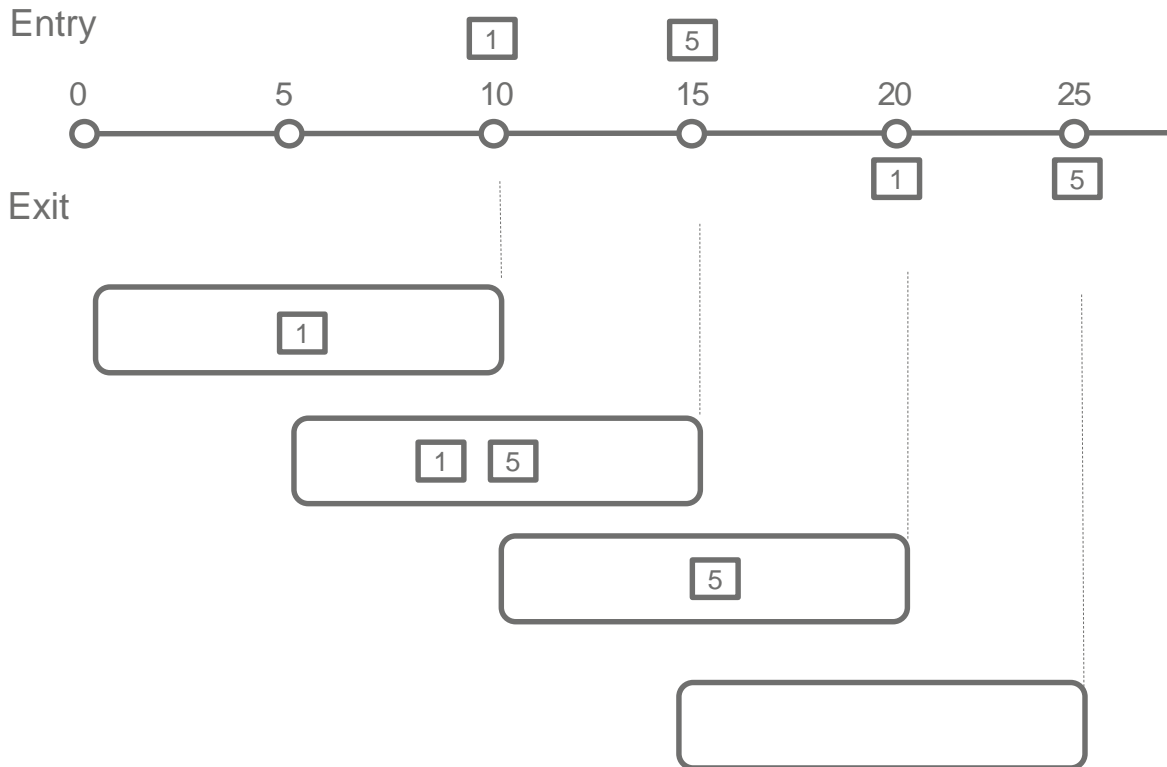


A 10-second Hopping Window with a 5-second "Hop"



Report every 5 seconds the total weight of cars that entered each toll both in the past 10 seconds

```
SELECT TollId, SUM(VehicleWeight) AS  
TotalWeight  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TollId, HoppingWindow(second, 10 , 5)
```




Report all toll booths which have served more than 3 vehicle in the last 10 seconds

```
SELECT TollId, Count(*)  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TollId, SlidingWindow(second, 10)  
HAVING COUNT(*) > 3
```

# Joining multiple streams

Calculate the time required for each car to pass the toll  
(in maximum 15 minutes)



```
SELECT EN.LicensePlate, DATEDIFF(minute, EN.EntryTime, EX.ExitTime) AS TotalTime
FROM EntryStream EN TIMESTAMP BY EntryTime
JOIN ExitStream EX TIMESTAMP BY ExitTime
  ON EN.TollId = EX.TollId AND EN.LicensePlate = EX.LicensePlate
  AND DATEDIFF(minute, EN, EX) BETWEEN 0 AND 15
```

# Detecting absence of events



Report all cars that did not pass the toll booth within 5 minutes



```
SELECT EN.LicensePlate, EN.TollId FROM EntryStream EN TIMESTAMP BY EntryTime
LEFT OUTER JOIN ExitStream EX TIMESTAMP BY ExitTime
  ON EN.TollId = EX.TollId AND EN.LicensePlate = EX.LicensePlate
  AND DATEDIFF(minute, EN, EX) BETWEEN 0 AND 5
WHERE EX.ExitTime IS NULL
```


Seamless correlation of event streams with reference data

Static or slowly-changing data stored in blobs

Scanned for changes on a settable cadence

**JOIN (INNER or LEFT OUTER)** between streams and reference data sources

Reference data appears like another input in the query



```
SELECT myRefData.Name, myStream.Value
FROM myStream
JOIN myRefData
    ON myStream.myKey = myRefData.myKey
```





- **Apache Spark:** Fast and general engine for large-scale data processing
- **Apache Storm:** Distributed real-time computation system to reliably process unbounded streams of data
- **Apache Flink:** Scalable stream and batch processing
- **Apache Ignite:** In-memory computing platform

## Commercial:

- Tibco Event Stream Processing
- IBM Streams
- Microsoft StreamInsight/ Azure Stream Analytics
- Amazon AWS Kinesis Analytics